

TimeSeriesAnalysis_TaoJin.R

my_macbook

Wed Oct 24 11:47:11 2018

```
library(knitr)
opts_chunk$set(message = FALSE, warning = FALSE, cache = TRUE, cache.lazy = FALSE)
options(width = 100, dplyr.width = 100)
library(ggplot2)
theme_set(theme_light())

library(quantmod)
```

```
## Loading required package: xts
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
## Loading required package: TTR
```

```
## Version 0.4-0 included new data defaults. See ?getSymbols.
```

```
library(forecast)
library(TTR)
library(tseries)

startDate = as.Date ("2010-06-29")
endDate = as.Date("2018-06-15")

# Get the Telsa Stock information from Yahoo

getSymbols("TSLA",src="yahoo",from = startDate, to = endDate)
```

```
## 'getSymbols' currently uses auto.assign=TRUE by default, but will
## use auto.assign=FALSE in 0.5-0. You will still be able to use
## 'loadSymbols' to automatically load data. getOption("getSymbols.env")
## and getOption("getSymbols.auto.assign") will still be checked for
## alternate defaults.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.
```

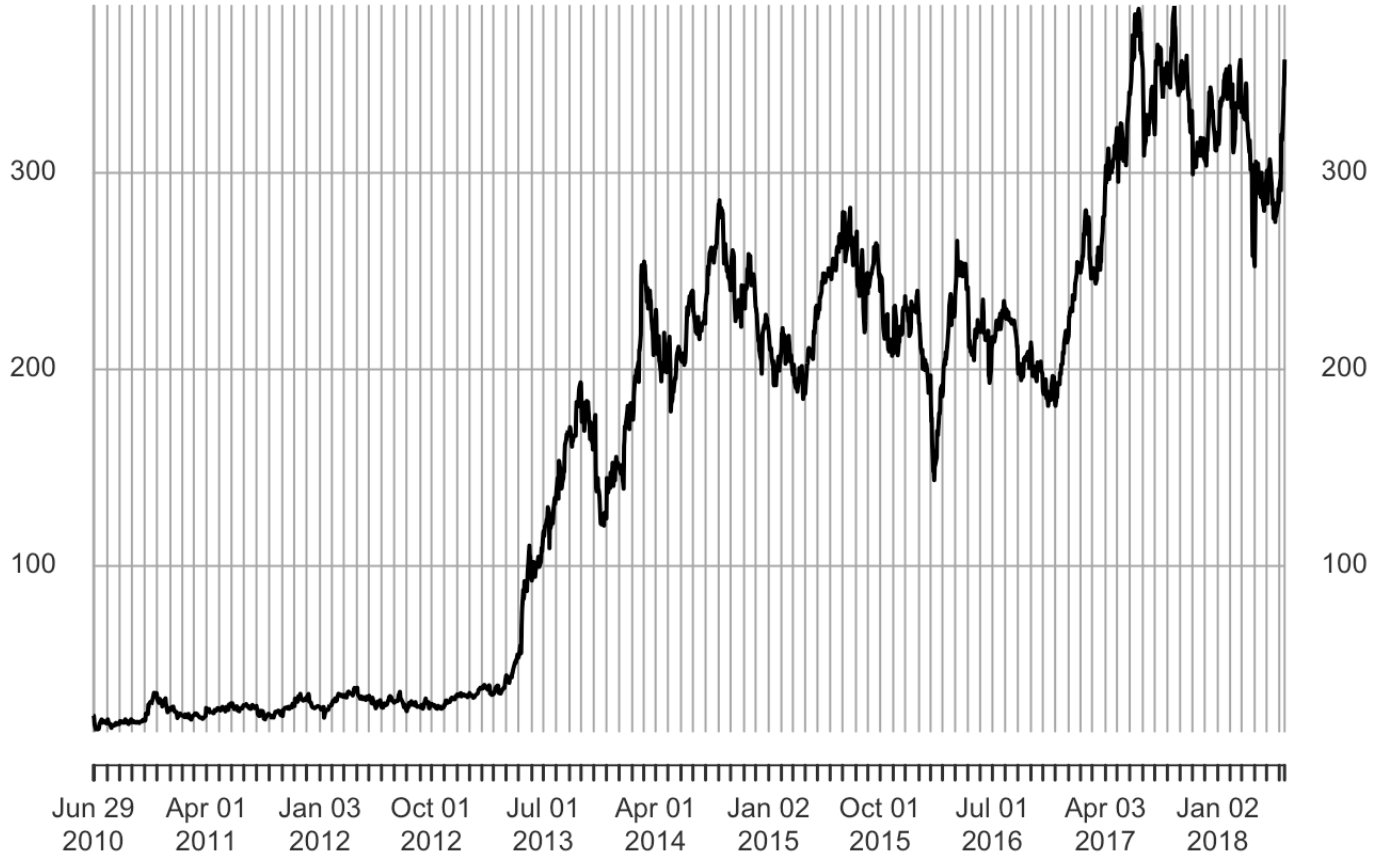
```
##
## WARNING: There have been significant changes to Yahoo Finance data.
## Please see the Warning section of '?getSymbols.yahoo' for details.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.yahoo.warning"=FALSE).
```

```
## [1] "TSLA"
```

```
plot(TSLA$TSLA.Close)
```


TSLA\$TSLA.Close

2010-06-29 / 2018-06-14



```
TeslaStockData <-ts(TSLA[,4], start = c(2010,6,29), end = c(2018,6,15), frequency = 365)
```


```
plot.ts(TeslaStockData)
```




*#We need to transform the time series in order to get a transformed time series that can be described using an additive
#model. transform the time series by calculating the natural log of the original data
:*

```
logTeslaStockData <- log(TeslaStockData)
```


```
plot.ts(logTeslaStockData)
```



```
#Decompose Time Series  
#Decomposing a time series means separating it into its constituent components, which  
are usually a trend component  
#and an irregular component, and if it is a seasonal time series, a seasonal componen  
t.  
# use the "SMA()" function to smooth time series data.  
  
logTeslaStockDataSMA8 <- SMA(logTeslaStockData, n=8)  
plot.ts(logTeslaStockDataSMA8)
```




```
logTeslaStockDataSMA40 <- SMA(logTeslaStockData, n=40)
plot.ts(logTeslaStockDataSMA40)
```



```
#Forecasts using Holt's Exponential Smoothing  
#Smoothing is controlled by two parameters, alpha, for the estimate of the level at t  
he current time point, and beta for  
#the estimate of the slope b of the trend component at the current time point.with si  
mple exponential smoothing, the paramters alpha  
#and beta have values between 0 and 1, and values that are close to 0 mean that littl  
e weight is placed on the most  
#recent observations when making forecasts of future values.
```


```
logTeslaStockDataForecast <- HoltWinters(logTeslaStockData,beta =FALSE, gamma=FALSE)  
plot(logTeslaStockDataForecast$fitted)
```

logTeslaStockDataForecast\$fitted



```
plot(logTeslaStockDataForecast)
```


Holt-Winters filtering




```
logTeslaStockDataForecast$SSE
```

```
## [1] 10.64326
```

```
# Forecast the next one year's price  
logTeslaStockDataForecast2 <- forecast(logTeslaStockDataForecast, h = 365)  
plot(logTeslaStockDataForecast2)
```

Forecasts from HoltWinters




*# in other words, if there are correlations between forecast errors for
successive predictions, it is likely that the simple exponential smoothing forecast
s could be improved upon by
another forecasting technique.*

*#We can calculate a correlogram of the forecast errors using the "acf()" function in
R. To specify the maximum lag
#that we want to look at, we use the "lag.max" parameter in acf().*

```
plot(acf(logTeslaStockDataForecast2$residuals, lag.max=40, na.action = na.pass))
```


Series logTeslaStockDataForecast2\$residuals



```
Box.test(logTeslaStockDataForecast2$residuals, lag=40, type="Ljung-Box")
```

```
##  
## Box-Ljung test  
##  
## data: logTeslaStockDataForecast2$residuals  
## X-squared = 31.42, df = 40, p-value = 0.832
```

```
plot.ts(logTeslaStockDataForecast2$residuals)
```




#To check whether the forecast errors are normally distributed with mean zero, we can plot a histogram of the forecast errors, with an overlaid normal curve that has mean zero and the same standard deviation as the distribution of forecast errors. To do this, we can define an R function "plotForecastErrors()", below:

```
plotForecastErrors <- function(forecasterrors)
{
  # make a histogram of the forecast errors:
  mybinsize <- IQR(forecasterrors)/4
  mysd <- sd(forecasterrors)
  mymin <- min(forecasterrors) - mysd*5
  mymax <- max(forecasterrors) + mysd*3
  # generate normally distributed data with mean 0 and standard deviation mysd
  mynorm <- rnorm(10000, mean=0, sd=mysd)
  mymin2 <- min(mynorm)
  mymax2 <- max(mynorm)
  if (mymin2 < mymin) { mymin <- mymin2 }
  if (mymax2 > mymax) { mymax <- mymax2 }
  # make a red histogram of the forecast errors, with the normally distributed data overlaid:
  mybins <- seq(mymin, mymax, mybinsize)
  hist(forecasterrors, col="red", freq=FALSE, breaks=mybins)
  # freq=FALSE ensures the area under the histogram = 1
  # generate normally distributed data with mean 0 and standard deviation mysd
  myhist <- hist(mynorm, plot=FALSE, breaks=mybins)
  # plot the normal curve as a blue line on top of the histogram of forecast errors:
  points(myhist$mids, myhist$density, type="l", col="blue", lwd=2)
}

logTeslaStockDataForecast2$residuals <- tsclean(logTeslaStockDataForecast2$residuals)

plotForecastErrors(logTeslaStockDataForecast2$residuals)
```

Histogram of forecasterrors



```
# it is plausible that the forecast errors are normally distributed with mean zero. The time plot of forecast errors shows
# that the forecast errors have roughly constant variance over time.
# The histogram of forecast errors show that it is plausible that the forecast errors are normally distributed with mean
# zero and constant variance.
```

```
# ARIMA Model
# While exponential smoothing methods do not make any assumptions about correlations between successive values
# of the time series, in some cases you can make a better predictive model by taking correlations in the data into
# account. Autoregressive Integrated Moving Average (ARIMA) models include an explicit statistical model for the
# irregular component of a time series, that allows for non-zero autocorrelations in the irregular component.

# Stationarized the Time Series
adf.test(logTeslaStockData)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: logTeslaStockData
## Dickey-Fuller = -1.6737, Lag order = 14, p-value = 0.7165
## alternative hypothesis: stationary
```


```
logTeslaStockData_Diff1 <- diff(logTeslaStockData, differences = 1)

adf.test(logTeslaStockData_Diff1)
```

```
## Warning in adf.test(logTeslaStockData_Diff1): p-value smaller than printed p-value
```


```
##
## Augmented Dickey-Fuller Test
##
## data: logTeslaStockData_Diff1
## Dickey-Fuller = -14.115, Lag order = 14, p-value = 0.01
## alternative hypothesis: stationary
```

```
plot(logTeslaStockData_Diff1)
```



```
acf(logTeslaStockData_Diff1, lag.max = 40)
```


Series logTeslaStockData_Diff1




```
acf(logTeslaStockData_Diff1, lag.max = 40, plot = FALSE)
```

```
##  
## Autocorrelations of series 'logTeslaStockData_Diff1', by lag  
##  
## 0.00000 0.00274 0.00548 0.00822 0.01096 0.01370 0.01644 0.01918 0.02192 0.02466 0.  
02740 0.03014  
## 1.000 -0.003 0.006 0.022 0.036 0.005 -0.034 -0.027 0.015 -0.029 -  
0.023 0.020  
## 0.03288 0.03562 0.03836 0.04110 0.04384 0.04658 0.04932 0.05205 0.05479 0.05753 0.  
06027 0.06301  
## -0.018 -0.021 0.021 0.011 -0.001 0.016 -0.002 0.015 0.002 0.000  
0.017 -0.007  
## 0.06575 0.06849 0.07123 0.07397 0.07671 0.07945 0.08219 0.08493 0.08767 0.09041 0.  
09315 0.09589  
## -0.001 0.011 0.005 0.003 0.003 0.004 0.032 -0.007 -0.017 0.006 -  
0.019 -0.007  
## 0.09863 0.10137 0.10411 0.10685 0.10959  
## -0.006 0.013 -0.011 0.023 -0.004
```

```
pacf(logTeslaStockData_Diff1, lag.max = 40)
```

Series logTeslaStockData_Diff1




```
pacf(logTeslaStockData_Diff1, lag.max = 40, plot = FALSE)
```

```
##
## Partial autocorrelations of series 'logTeslaStockData_Diff1', by lag
##
## 0.00274 0.00548 0.00822 0.01096 0.01370 0.01644 0.01918 0.02192 0.02466 0.02740 0.
03014 0.03288
## -0.003 0.006 0.022 0.036 0.005 -0.035 -0.029 0.013 -0.027 -0.020
0.022 -0.018
## 0.03562 0.03836 0.04110 0.04384 0.04658 0.04932 0.05205 0.05479 0.05753 0.06027 0.
06301 0.06575
## -0.020 0.022 0.010 -0.002 0.018 -0.003 0.010 0.003 0.001 0.014 -
0.006 0.001
## 0.06849 0.07123 0.07397 0.07671 0.07945 0.08219 0.08493 0.08767 0.09041 0.09315 0.
09589 0.09863
## 0.010 0.006 0.005 0.003 0.004 0.031 -0.006 -0.017 0.004 -0.020 -
0.005 -0.004
## 0.10137 0.10411 0.10685 0.10959
## 0.015 -0.011 0.025 -0.004
```


```
TeslaARIMA <- auto.arima(logTeslaStockData_Diff1)

TeslaARIMA_Forecast <-forecast(TeslaARIMA, h = 365)
plot(TeslaARIMA_Forecast)
```

Forecasts from ARIMA(0,0,0) with zero mean

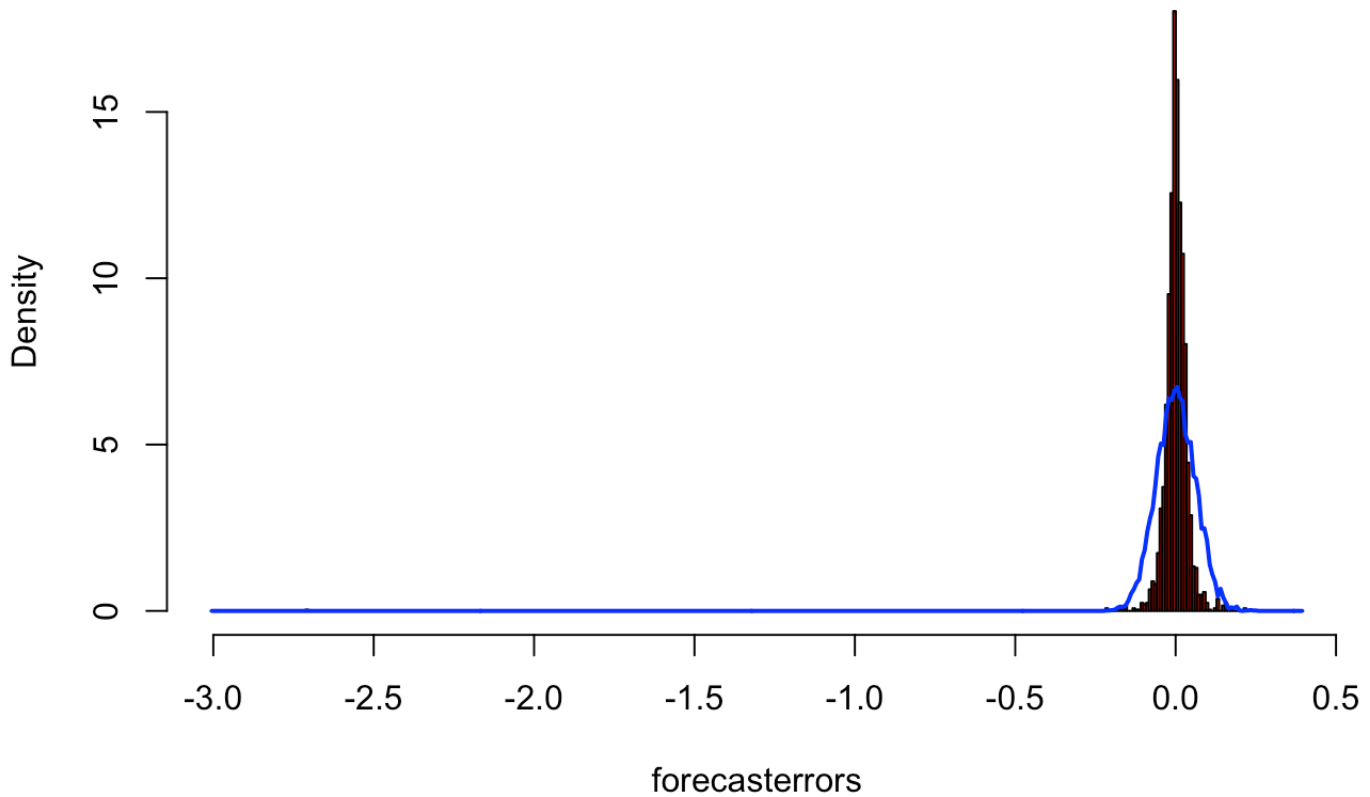


```
plot(TeslaARIMA_Forecast$residuals)
```



```
plotForecastErrors(TeslaARIMA_Forecast$residuals)
```

Histogram of forecasterrors




```
# White noise problem, choose the ARIMA (0,1,0) instead of auto.arima


TeslaARIMA1 <- arima(logTeslaStockData_Diff1, order = c(0,1,0))
TeslaARIMA1_Forecast <- forecast(TeslaARIMA1, h = 365)

plot(TeslaARIMA1_Forecast)
```

Forecasts from ARIMA(0,1,0)

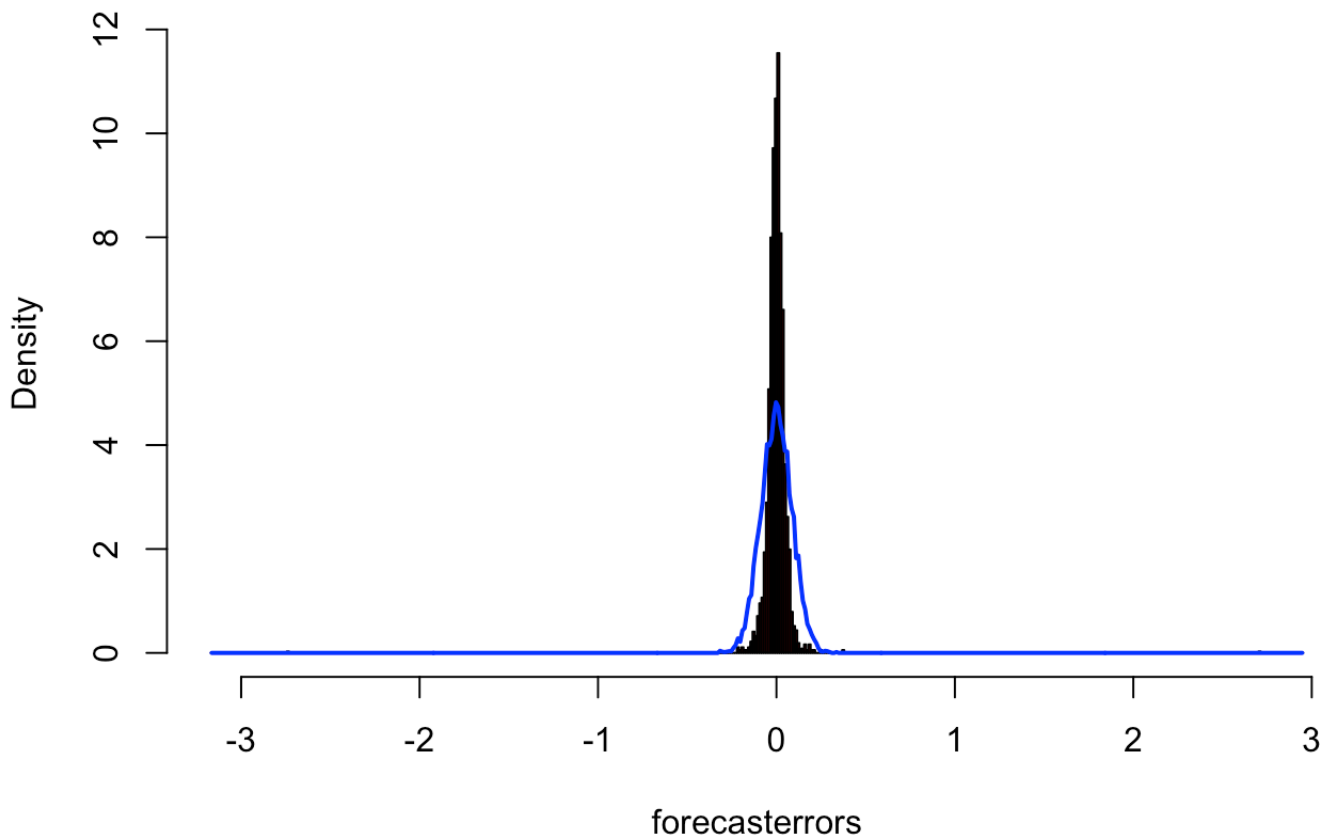


```
plot(TeslaARIMA1_Forecast$residuals)
```

```
plotForecastErrors(TeslaARIMA1_Forecast$residuals)
```

Histogram of forecasterrors



```
library(sarima)
```

```
## Loading required package: stats4
```

```
library(stats4)  
# whiteNoiseTest(TeslaARIMA1_Forecast,h0= iid)
```